Intelligence for the physical world.

AR without glasses. Projected light that guides real tasks, step by step.

The Light Company is building intelligence that understands physical work, places instructions on real objects and surfaces, verifies progress, and guides a task to completion. Robotics is the first wedge. The larger platform brings useful AR into assembly, repair, training, and everyday work.

10x

EVALUATION EFFICIENCY

By run count: 300 to 30

92.5%

LOWER VARIABILITY

40% to 3% in 30 days

10 paid units | 3 paying customers | 20 stations running in labs
Enact: 5 after one demo | New Theory: 4 | Trossen Robotics: 1.

The problem: the physical input changes between tests

A robot team needs to know whether its new policy improved. But someone must reset the objects before every trial. Even at a fixed station, a different position or orientation can change the result. That noise can distort model comparisons, debugging, and decisions about what data to collect.

What Prism solves: a repeatable, verified starting condition

Prism projects the target placement into the workspace. An operator resets the scene. Cameras verify the placement against the target. The light turns off before the robot runs. The result can then be linked to a known starting state. Teams can repeat the same test or vary conditions deliberately.

Enact evaluations fell from 300-400 runs to 30 (10x). New Theory's variability fell from 40% to 3% in 30 days. Prism scores every run against the team's rubric and records the reason.

Replace guesswork with engineering

Enact (YC S26) bought five units after one demo for consistent initial conditions. The planned software would map failures and undersampled states: places where a large dataset may still lack useful examples. Teams could collect targeted data, retrain, and retest the same conditions. The goal is a measured learning loop, with broader AR applications in assembly, repair, and training.

One station sale. Software on every run.

A station sells for $1,500 at an 80% margin. Every evaluation after that runs through Prism at $1, every engineer works in the platform, and tooling and enterprise contracts sit on top. Labs re-test every time the model changes, so the software recurs. **250 labs at 20 stations each: $187M a year in software.** Applied Intuition, the same layer for self-driving cars, is valued at $15 billion.

Raising our seed to ship 100 stations into labs, make the software self-serve, and publish the first public benchmark.

Sources: company records; YC / Enact; CB Insights; Bloomberg.

Jonathan Solomon | hi@lightcompany.ai | lightcompany.ai